

Отчет к ДЗ Репликация и удаление

1. Взять любой демонстрационный DATASET, не обязательно полный набор данных:
<https://clickhouse.com/docs/en/getting-started/example-datasets>
2. Конвертировать таблицу в реплицируемую, используя макрос replica
3. Добавить 2 реплики
4. Отдать результаты запросов как 2 файла

```
SELECT
```

```
getMacro('replica'),
```

```
*
```

```
FROM remote('разделенный запятыми список реплик',system.parts)
```

```
FORMAT JSONEachRow;
```

```
SELECT * FROM system.replicas FORMAT JSONEachRow;
```

5. Добавить/выбрать колонку с типом Date в таблице, добавить TTL на таблицу «хранить последние 7 дней». Предоставить результат запроса «SHOW CREATE TABLE таблица» на проверку.

1. Создаем таблицу из тестовых датасетов и наполняем ее

```
53d58e10c49a :) CREATE TABLE uk_price_paid
(
  price UInt32,
  date Date,
  postcode1 LowCardinality(String),
  postcode2 LowCardinality(String),
  type Enum8('terraced' = 1, 'semi-detached' = 2, 'detached' = 3, 'flat' = 4, 'other' = 0),
  is_new UInt8,
  duration Enum8('freehold' = 1, 'leasehold' = 2, 'unknown' = 0),
  addr1 String,
  addr2 String,
  street LowCardinality(String),
  locality LowCardinality(String),
  town LowCardinality(String),
  district LowCardinality(String),
  county LowCardinality(String)
)
ENGINE = MergeTree
ORDER BY (postcode1, postcode2, addr1, addr2);

CREATE TABLE uk_price_paid
(
  `price` UInt32,
  `date` Date,
  `postcode1` LowCardinality(String),
  `postcode2` LowCardinality(String),
  `type` Enum8('terraced' = 1, 'semi-detached' = 2, 'detached' = 3, 'flat' = 4, 'other' = 0),
  `is_new` UInt8,
  `duration` Enum8('freehold' = 1, 'leasehold' = 2, 'unknown' = 0),
  `addr1` String,
  `addr2` String,
  `street` LowCardinality(String),
  `locality` LowCardinality(String),
  `town` LowCardinality(String),
  `district` LowCardinality(String),
  `county` LowCardinality(String)
)
ENGINE = MergeTree
ORDER BY (postcode1, postcode2, addr1, addr2)

Query id: b772afc5-4cd6-4af7-94f2-243f63c9d8ea

OK.

0 rows in set. Elapsed: 0.010 sec.

INSERT INTO uk_price_paid
SETTINGS max_http_get_redirects = 10
WITH splitByChar(' ', postcode) AS p
SELECT
  toUInt32(price_string) AS price,
  parseDateTimeBestEffortUS(time) AS date,
  p[1] AS postcode1,
  p[2] AS postcode2,
  transform(a, ['T', 'S', 'D', 'F', 'O'], ['terraced', 'semi-detached', 'detached', 'flat', 'other']) AS type,
  b = 'Y' AS is_new,
  transform(c, ['F', 'L', 'U'], ['freehold', 'leasehold', 'unknown']) AS duration,
  addr1,
  addr2,
  street,
  locality,
  town,
  district,
  county
FROM url('http://prod.publicdata.landregistry.gov.uk/s3-website-eu-west-1.amazonaws.com/pp-complete.csv', 'CSV', 'uuid_string String,\n price_s\n\n locality String,\n town String,\n district String,\n county String,\n d String,\n e String')
SETTINGS max_http_get_redirects = 10

Query id: 4a51dc16-27c8-4c9d-a8db-67ab1ff74fbf

Progress: 1.85 million rows, 325.06 MB (211.03 thousand rows/s., 37.02 MB/s.)
Press the space key to toggle the display of the progress table.
```

```
SELECT count(*)
FROM uk_price_paid
```

Query id: 9ed57d8f-dcae-447a-a33b-fdbf0165dd0a

```
count()
1. 29946800 -- 29.95 million
```

1 row in set. Elapsed: 0.003 sec.

53d58e10c49a :) █

2. Конвертируем таблицу в реплицируемую и аттачим в нее партиции из обычной таблицы:

```
CREATE TABLE uk_price_paid_repl AS uk_price_paid  
ENGINE = ReplicatedMergeTree('/clickhouse/{cluster}/{table}', '{replica}')
```

Query id: 5aadcee0-4812-4f33-97fd-84469700bf67

Ok.

0 rows in set. Elapsed: 0.075 sec.

```
53d58e10c49a :) ALTER TABLE uk_price_paid_repl ATTACH PARTITION ID 'all' FROM uk_price_paid;
```

```
ALTER TABLE uk_price_paid_repl  
(ATTACH PARTITION ID 'all' FROM uk_price_paid)
```

Query id: 968c29e1-7ad0-4d1d-a4da-6411e0875615

Ok.

0 rows in set. Elapsed: 0.071 sec.

```
53d58e10c49a :) select count(*) from uk_price_paid_repl;
```

```
SELECT count(*)  
FROM uk_price_paid_repl
```

Query id: a7f08fa8-443f-4842-b551-9c6609cf967d

```
1.  count(*)  
    29946800  -- 29.95 million
```

1 row in set. Elapsed: 0.003 sec.

3. Из той же ноды создаю таблицы на двух других нодах и проверяем что туда данные реплицировались:

```
CREATE TABLE IF NOT EXISTS uk_price_paid_repl ON CLUSTER '{cluster}'
(
    `price` UInt32,
    `date` Date,
    `postcode1` LowCardinality(String),
    `postcode2` LowCardinality(String),
    `type` Enum8('terraced' = 1, 'semi-detached' = 2, 'detached' = 3, 'flat' = 4, 'other' = 0),
    `is_new` UInt8,
    `duration` Enum8('freehold' = 1, 'leasehold' = 2, 'unknown' = 0),
    `addr1` String,
    `addr2` String,
    `street` LowCardinality(String),
    `locality` LowCardinality(String),
    `town` LowCardinality(String),
    `district` LowCardinality(String),
    `county` LowCardinality(String)
)
ENGINE = ReplicatedMergeTree('/clickhouse/{cluster}/{table}', '{replica}')
ORDER BY (postcode1, postcode2, addr1, addr2)
```

Query id: 122ec04b-4e8f-4ab4-aa98-dc1e8f6245fa

	host	port	status	error	num_hosts_remaining	num_hosts_active
1.	clickhouse1	9000	0		2	2
2.	clickhouse2	9000	0		1	0
3.	clickhouse3	9000	0		0	0

3 rows in set. Elapsed: 0.137 sec.

53d58e10c49a :)

c6866125ae81 :) show tables

SHOW TABLES

Query id: fc26d1f5-cfef-4bc7-a2f7-a43fffd0a082

	name
1.	example
2.	uk_price_paid_repl

2 rows in set. Elapsed: 0.003 sec.

c6866125ae81 :) select COUNT(*) from uk_price_paid_repl;

SELECT COUNT(*)
FROM uk_price_paid_repl

Query id: 8316d49c-d8ac-4106-9fa8-7711c93c7e28

	COUNT()	
1.	29946800	-- 29.95 million

1 row in set. Elapsed: 0.003 sec.

c6866125ae81 :)

SHOW TABLES

Query id: 31e30913-52ec-4a27-88a0-74262c8e91db

	name
1.	example
2.	uk_price_paid_repl

2 rows in set. Elapsed: 0.002 sec.

deffb0ff20ce :) select count(*) FROM uk_price_paid_repl;

SELECT count(*)
FROM uk_price_paid_repl

Query id: 14a0879e-e9e9-40b8-85e6-28d416ecb640

	count()	
1.	29946800	-- 29.95 million

1 row in set. Elapsed: 0.002 sec.

deffb0ff20ce :)

4. Файлы как результаты запроса:



output_1.json output_2.json

5. Изменение таблицы. Добавлчем в нее TTL на последние 7 дней:

```
53d58e10c49a :) ALTER TABLE uk_price_paid_repl on cluster '{cluster}' MODIFY TTL date + INTERVAL 7 DAY;  
  
ALTER TABLE uk_price_paid_repl ON CLUSTER '{cluster}'  
  (MODIFY TTL date + toIntervalDay(7))
```

Query id: 10c67796-b8a9-4967-baf6-357d1fd051ca

	host	port	status	error	num_hosts_remaining	num_hosts_active
1.	clickhouse2	9000	0		2	0
2.	clickhouse3	9000	0		1	0
3.	clickhouse1	9000	0		0	0

3 rows in set. Elapsed: 8.601 sec.

53d58e10c49a :) █

```
SHOW CREATE TABLE uk_price_paid_repl  
FORMAT Native
```

Query id: 023b602d-c859-4266-a037-37232718be7d

The requested output format 'Native' is binary and could produce side-effects when output directly into the terminal.
If you want to output it into a file, use the "INTO OUTFILE" modifier in the query or redirect the output of the shell command.
Do you want to output it anyway? [y/N] Y

```
statementStringCREATE TABLE default.uk_price_paid_repl  
(  
  `price` UInt32,  
  `date` Date,  
  `postcode1` LowCardinality(String),  
  `postcode2` LowCardinality(String),  
  `type` Enum8('other' = 0, 'terraced' = 1, 'semi-detached' = 2, 'detached' = 3, 'flat' = 4),  
  `is_new` UInt8,  
  `duration` Enum8('unknown' = 0, 'freehold' = 1, 'leasehold' = 2),  
  `addr1` String,  
  `addr2` String,  
  `street` LowCardinality(String),  
  `locality` LowCardinality(String),  
  `town` LowCardinality(String),  
  `district` LowCardinality(String),  
  `county` LowCardinality(String)  
)  
ENGINE = ReplicatedMergeTree('/clickhouse/{cluster}/uk_price_paid_repl', '{replica}')  
ORDER BY (postcode1, postcode2, addr1, addr2)  
TTL date + toIntervalDay(7)  
SETTINGS index_granularity = 8192
```

1 row in set. Elapsed: 0.002 sec.

c6866125ae81 :)

```
SHOW CREATE TABLE uk_price_paid_repl
```

Query id: fede78e2-80cc-4b2c-81d0-0dc3b99822df

```
statement  
1. CREATE TABLE default.uk_price_paid_repl  
  (  
    `price` UInt32,  
    `date` Date,  
    `postcode1` LowCardinality(String),  
    `postcode2` LowCardinality(String),  
    `type` Enum8('other' = 0, 'terraced' = 1, 'semi-detached' = 2, 'detached' = 3, 'flat' = 4),  
    `is_new` UInt8,  
    `duration` Enum8('unknown' = 0, 'freehold' = 1, 'leasehold' = 2),  
    `addr1` String,  
    `addr2` String,  
    `street` LowCardinality(String),  
    `locality` LowCardinality(String),  
    `town` LowCardinality(String),  
    `district` LowCardinality(String),  
    `county` LowCardinality(String)  
  )  
ENGINE = ReplicatedMergeTree('/clickhouse/{cluster}/uk_price_paid_repl', '{replica}')  
ORDER BY (postcode1, postcode2, addr1, addr2)  
TTL date + toIntervalDay(7)  
SETTINGS index_granularity = 8192
```

1 row in set. Elapsed: 0.001 sec.

c6866125ae81 :)